

What is Database:

A database is a structured way to store, organize, and manage data so it can be easily accessed, updated, and analyzed. Think of it as a digital filing system where data is stored in tables, rows, and columns, like a spreadsheet. It helps applications and users retrieve or manipulate data efficiently.

Examples: Storing customer details for an e-commerce site or keeping track of bank transactions.

Types of Databases

1. **Relational Databases (SQL):**
 - Stores data in tables with rows and columns.
 - Uses Structured Query Language (SQL) for managing data.
 - Examples: MySQL, PostgreSQL, Oracle.
2. **NoSQL Databases:**
 - Stores unstructured or semi-structured data (JSON, key-value pairs, documents).
 - Flexible schema for dynamic data.
 - Examples: MongoDB, DynamoDB.

Why Use a Database?

1. **Efficiency:**
 - Faster access and retrieval of data compared to files or spreadsheets.
2. **Scalability:**
 - Handles large volumes of data seamlessly.
3. **Data Integrity:**
 - Ensures accuracy and consistency with rules and relationships.
4. **Multi-User Access:**
 - Supports concurrent access by multiple users.
5. **Security:**
 - Provides advanced security controls to protect sensitive data.
6. **Backup and Recovery:**
 - Easy to back up and recover in case of failure.
7. **Querying and Analysis:**
 - Simplifies complex data searches and analytics with languages like SQL.

Databases are more robust, reliable, and efficient for structured, large, or growing datasets.

AWS RDS (Relational Database Service) Overview:

What is AWS RDS?

- AWS RDS is a **managed database service** that simplifies the setup, operation, and scaling of relational databases in the cloud.
 - Supports various database engines:
 - **Amazon Aurora** (MySQL and PostgreSQL compatible)
 - **MySQL**
 - **PostgreSQL**
 - **MariaDB**
 - **Oracle**
 - **Microsoft SQL Server**
-

Why Use AWS RDS Instead of Setting Up Databases on EC2?

1. **Ease of Management:**
 - Fully managed service: No manual installation, patching, or backups.
 - Automated backups, snapshots, and point-in-time recovery.
2. **Scalability:**
 - Vertical scaling: Change instance types for more resources.
 - Horizontal scaling: Read replicas for improved read performance.
3. **High Availability:**
 - Multi-AZ deployments for automatic failover.
 - Automatically manages standby and primary instances.
4. **Performance:**
 - Optimized storage options like **General Purpose SSD**, **Provisioned IOPS SSD**.
 - Amazon Aurora offers superior performance compared to traditional MySQL/PostgreSQL.
5. **Security:**
 - Built-in encryption for data at rest (using AWS KMS) and in transit (using SSL/TLS).
 - Integration with AWS IAM for fine-grained access control.
 - Network isolation using VPC.
6. **Cost-Effectiveness:**
 - Pay-as-you-go model.
 - Saves costs by reducing administrative overhead.
 - No need to provision and maintain EC2 instances for database hosting.

7. **Monitoring and Metrics:**
 - Built-in monitoring via **Amazon CloudWatch**.
 - Database performance insights for query optimization.
-

Key Features of AWS RDS

1. **Database Engines:**
 - Choose from six popular relational database engines.
 2. **Automated Backups:**
 - Daily snapshots and transaction logs enable recovery to any point within the retention period.
 3. **Multi-AZ Deployments:**
 - Increases database availability with automatic failover.
 4. **Read Replicas:**
 - Improves read performance by distributing read queries to replicas.
 5. **Storage Options:**
 - **General Purpose SSD:** Balanced performance and cost.
 - **Provisioned IOPS SSD:** High-performance, low-latency workloads.
 - **Magnetic Storage:** Cost-effective for infrequent access.
 6. **Maintenance:**
 - Automatic patching of the database engine.
 7. **Security:**
 - Database encryption using AWS Key Management Service (KMS).
 - Access control via IAM and security groups.
 8. **Support for VPC:**
 - Database instances can run in a Virtual Private Cloud for network isolation.
 9. **Monitoring:**
 - CloudWatch metrics, alarms, and logs for detailed performance insights.
-

When to Use AWS RDS?

- **Production workloads** requiring high availability and durability.
 - Applications needing managed database solutions to minimize administrative overhead.
 - Workloads needing scalability for growing traffic.
 - Scenarios requiring encryption for data compliance (e.g., healthcare, finance).
 - Applications with heavy read workloads (use Read Replicas).
-

Key Differences Between RDS and Database on EC2

Feature	AWS RDS	Database on EC2
---------	---------	-----------------

Management	Fully managed by AWS	Fully managed by you
Setup	Simple via console or CLI	Manual installation/configuration
Scalability	Auto-scaling supported	Manual scaling
Backups	Automated backups	Manual backups
High Availability	Built-in Multi-AZ support	Requires custom setup
Performance Tuning	Pre-configured optimizations	Fully manual
Cost	Pay for usage	Pay for EC2, storage, and management

Considerations When Using RDS

1. **Limitations:**
 - No OS-level access (unlike EC2).
 - Limited customization of database parameters compared to on-prem databases.
 2. **Use Cases:**
 - **Transactional systems:** E-commerce, banking.
 - **Read-heavy applications:** Blogs, analytics.
 - **SaaS platforms** needing managed databases.
 3. **Advanced Features:**
 - **Aurora Global Database:** Globally distributed workloads.
 - **RDS Proxy:** Connection pooling to improve app performance and availability.
-

RDS Best Practices

- Use **Multi-AZ deployments** for critical applications.
- Enable **automated backups** and **retention periods** for disaster recovery.
- Monitor **CloudWatch metrics** to optimize performance.
- Secure access using **IAM roles**, **security groups**, and **SSL/TLS**.
- Use **Read Replicas** for scaling reads without impacting the primary database.

Database setup In EC2:

```
$ sudo apt install mysql-server
```

```
$ sudo systemctl start mysql.service
```

```
$ mysql -u root -p
```

Create a DB:

```
$ CREATE DATABASE my_database;
```

```
$ USE my_database;
```

```
$ CREATE TABLE users ( user_id INT AUTO_INCREMENT PRIMARY KEY, first_name  
VARCHAR(50), last_name VARCHAR(50), email VARCHAR(100), date_of_birth DATE );
```

```
$ INSERT INTO users (first_name, last_name, email, date_of_birth) VALUES ('John', 'Doe',  
'john.doe@example.com', '1990-01-01');
```

```
$ SELECT * FROM users;
```

```
$ INSERT INTO users (first_name, last_name, email, date_of_birth) VALUES ('John', 'Doe',  
'john.doe@example.com', '1990-01-01');
```

```
$ EXIT;
```

Connect to RDS from an EC2:

```
$ sudo apt-get install mysql-client
```

```
$ mysql -h database-1.c9y608sqq15u.us-east-1.rds.amazonaws.com -u admin -p
```

—X—

Laxmikanat Giri

